

# ESA CSOC/ESOC — Rsyslog Load-Balancer on RHEL 9 (TCP/TLS Round-Robin)

---

**Author:** Gaetan De Dobbeleer (Logpoint)

**Date:** 12 Aug 2025

**Scope:** Implement a robust rsyslog relay on RHEL 9 that load-balances outgoing syslog traffic to multiple Logpoint collectors/backends over **TCP/TLS** using the **native omfwd target pool** feature. Ports are **fixed by Logpoint: UDP 514, TCP 514 (clear), TCP 6514 (TLS)**.

---

## 1) Executive summary

- **OS:** RHEL 9 (standard build; no containers).
- **Rsyslog: v8.2502+ (recommended)** from the Adiscon RHEL 9 repo.
  - Rationale: native **load-balancing in omfwd** requires **≥ 8.2408**; v8.2502 additionally fixes a TLS handshake blocking issue and brings reliability/perf improvements.
- **Transport & fixed ports (Logpoint): UDP 514, TCP 514 (clear), TCP 6514 (TLS).**
- **Ingress on relay:** listen on UDP 514, TCP 514, and TCP 6514 (TLS).
- **Egress from relay: preferred** over **TCP 6514 (TLS)** with round-robin LB; optional cleartext egress over **TCP 514** provided (disabled by default).
- **Buffering & failover:** disk-assisted queues with automatic retry; if one target is down the pool rotates; if **all** targets are down, events are **buffered on disk** and automatically replayed when a backend returns.
- **Observability:** **impstats** every 60s, JSON/CEE format; actionable counters for queue depth, submits, failures, and suspension state.
- **Security:** SELinux **Enforcing**; firewalld opened only for 514/udp, 514/tcp, 6514/tcp; TLS defaults set (anonymous auth by default, can switch to mTLS later).

**Outcome:** transparent round-robin distribution of events across multiple Logpoint backends with automatic removal/retry of down targets, safe buffering when none are reachable, and clean recovery with no manual intervention.

---

## 2) Why this rsyslog version? (Design decision)

- The **native load-balancing** (array syntax **target=["ip1","ip2",...]** in **omfwd**) landed in **8.2408.0**.
  - RHEL 9 stock repos carry **8.2310.x**, which **does not** include that feature.
  - Installing from the **Adiscon RHEL 9 repo** provides **current v8** quickly and safely.
  - **Recommendation: 8.2502+** (released Feb 2025) — includes a **TLS handshake fix** preventing rare stalls in **omfwd** under load.
  - Acceptable “scheduled-stable” alternative if change control is strict: **8.2412.0** (still includes LB).
- 

## 3) Prerequisites

- RHEL 9 server (x86\_64), root access.

- Network reachability to Logpoint backends.
- **Fixed Logpoint ports: UDP 514, TCP 514 (clear), TCP 6514 (TLS).**
- TLS files (if TLS enabled):
  - `/etc/rsyslog.d/tls/ca.crt`
  - `/etc/rsyslog.d/tls/server.crt`
  - `/etc/rsyslog.d/tls/server.key` (0600)

**Backends:** `<BACKEND_1>`, `<BACKEND_2>`

---

## 3bis) Customization — Replace placeholders with ESA production targets

Before going live, **replace the placeholders** with the actual ESA Logpoint collector IPs or hostnames:

- `<BACKEND_1>` — first Logpoint backend (collector)
- `<BACKEND_2>` — second Logpoint backend (collector)

Quick example with `sed` (replace with your real values):

```
sudo sed -i 's/<BACKEND_1>/192.0.2.10/' /etc/rsyslog.d/10-esa-lb.conf
sudo sed -i 's/<BACKEND_2>/192.0.2.11/' /etc/rsyslog.d/10-esa-lb.conf
sudo systemctl restart rsyslog
```

Connectivity test (optional):

```
# TLS 6514
timeout 3 bash -c 'cat < /dev/null > /dev/tcp/<BACKEND_1>/6514' && echo "OK" ||
echo "FAIL"
timeout 3 bash -c 'cat < /dev/null > /dev/tcp/<BACKEND_2>/6514' && echo "OK" ||
echo "FAIL"
# Clear TCP 514 if you intend to enable it later
timeout 3 bash -c 'cat < /dev/null > /dev/tcp/<BACKEND_1>/514' && echo "OK" ||
echo "FAIL"
timeout 3 bash -c 'cat < /dev/null > /dev/tcp/<BACKEND_2>/514' && echo "OK" ||
echo "FAIL"
```

---

## 4) Installation on RHEL 9 (no Docker)

1. **Enable Adiscon repo & install** (v8-stable or daily-stable):

```
cd /etc/yum.repos.d/
sudo curl -O https://rpms.adiscon.com/v8-stable/rsyslog-rhel.repo
# (option) daily-stable for faster fixes
sudo curl -O https://rpms.adiscon.com/v8-stable-daily/rsyslog-daily-
rhel.repo
```

```
# ensure gpgcheck is enabled in the .repo file (security hygiene)
sudo sed -i 's/^gpgcheck=.*gpgcheck=1/' /etc/yum.repos.d/rsyslog-*.repo

sudo dnf clean all && sudo dnf makecache
sudo dnf install -y rsyslog rsyslog-gnutls
rsyslogd -v    # verify: aim for 8.2502+ (or ≥ 8.2412 at minimum)
```

## 2. Open firewalld (fixed ports):

```
sudo firewall-cmd --permanent --add-port=514/tcp
sudo firewall-cmd --permanent --add-port=514/udp
sudo firewall-cmd --permanent --add-port=6514/tcp
sudo firewall-cmd --reload
```

## 3. SELinux — keep **Enforcing**. If a non-default port is ever used, map it:

```
# example for TCP 30514 if used one day:
sudo semanage port -a -t syslogd_port_t -p tcp 30514
```

---

# 5) Configuration layout

## Files

- `/etc/rsyslog.conf` — core, inputs, TLS defaults, `impstats`.
- `/etc/rsyslog.d/10-esa-lb.conf` — outbound to Logpoint via LB + buffering + failover.
- `/etc/rsyslog.d/tls/` — TLS material.

### 5.1 `/etc/rsyslog.conf` (excerpt)

```
# Work directory for persistent queues (disk-assisted)
global(workDirectory="/var/spool/rsyslog")

# Inputs
module(load="imudp")
module(load="imtcp")
module(load="impstats" interval="60" format="cee") # stats/health

# TLS defaults (used by listener on 6514 and egress if enabled)
global(
    defaultNetstreamDriver="gtls"
    defaultNetstreamDriverCAFile="/etc/rsyslog.d/tls/ca.crt"
    defaultNetstreamDriverCertFile="/etc/rsyslog.d/tls/server.crt"
    defaultNetstreamDriverKeyFile="/etc/rsyslog.d/tls/server.key"
)

# Inputs from sources (fixed ports)
```

```

input(type="imudp" port="514" ruleset="to_logpoint")
input(type="imtcp" port="514" ruleset="to_logpoint")
input(type="imtcp" port="6514" StreamDriver.name="gtls" StreamDriver.mode="1"
StreamDriver.authmode="anon" ruleset="to_logpoint")

# Include drop-ins
$IncludeConfig /etc/rsyslog.d/*.conf

```

#### Notes:

- `omfwd` is builtin (no `module(load="omfwd")`).
- Template used later: `RSYSLOG_SyslogProtocol23Format` (RFC5424-like).

## 5.2 `/etc/rsyslog.d/10-esa-lb.conf` — **LB + buffering + failover**

```

# Keep impstats out of the forward stream (optional)
if ($syslogtag == 'impstats:') then {
    action(type="omfile" file="/var/log/rsyslog_stats.json")
    stop
}

# Main forwarding ruleset to Logpoint
ruleset(name="to_logpoint") {
    # --- Preferred: TLS 6514 with native round-robin load-balancing
    action(
        name="lp_tls_rr"
        type="omfwd" protocol="tcp"
        StreamDriver="gtls" StreamDriverMode="1" StreamDriverAuthMode="anon"
        target=["<BACKEND_1>", "<BACKEND_2>"] # round-robin pool
        port="6514"
        template="RSYSLOG_SyslogProtocol23Format"

        # --- Buffering (disk-assisted action queue)
        queue.type="LinkedList" # async, supports DA-queue
        queue.filename="q_logpoint_tls" # enables spooling under
workDirectory
        queue.maxdiskspace="10g" # cap for on-disk backlog
        queue.size="50000" # in-memory elements
        queue.highwatermark="40000"
        queue.lowwatermark="10000"
        queue.dequeuebatchsize="1024"
        queue.workerthreads="2"
        queue.saveonshutdown="on" # persist on planned reboot

        # --- Availability / retry policy
        action.resumeRetryCount="-1" # retry forever
        action.resumeInterval="30" # backoff between attempts (s)
    )

    # --- Optional (disabled): clear TCP 514 with round-robin load-balancing
    # action(

```

```
# name="lp_tcp_rr" type="omfwd" protocol="tcp"
# target=["<BACKEND_1>", "<BACKEND_2>"] port="514"
# template="RSYSLOG_SyslogProtocol23Format"
# queue.type="LinkedList" queue.filename="q_logpoint_tcp"
# queue.maxdiskspace="10g" queue.size="50000"
# queue.highwatermark="40000" queue.lowwatermark="10000"
# queue.dequeuebatchsize="1024" queue.workerthreads="2"
# queue.saveonshutdown="on"
# action.resumeRetryCount="-1" action.resumeInterval="30"
# )

# --- Local fallback (executes only if previous action is suspended)
action(
  name="local_fallback" type="omfile"
  file="/var/log/esa_fallback-buffer.log"
  execOnlyWhenPreviousIsSuspended="on"
)
}
```

### Behavior

- If one backend is **down**, **omfwd** **skips it** and sends to the next one in the pool.
- If **all** backends are down, the action becomes **suspended** and events are **buffered** (memory → disk) until a target returns.
- On recovery, rsyslog **automatically resumes** and **replays** the backlog in order.
- The **local fallback** only activates while the TLS action is suspended, ensuring you still retain local evidence.

## 6) Verification & smoke tests

```
# Syntax check
sudo rsyslogd -N1

# Enable & start
sudo systemctl enable --now rsyslog
sudo systemctl status rsyslog --no-pager

# Send a test event
logger -t ESA_SMOKE "rsyslog LB end-to-end OK"
```

### Operational checks

- Round-robin: temporarily stop one backend and verify events continue to the remaining target; restore and observe alternation.
- Buffering: stop **both** backends, generate traffic, confirm growth of spool files under `/var/spool/rsyslog/` and of the impstats queue metrics; restart backends and confirm automatic drain.

## 7) Monitoring / operations with **impstats**

We enabled **impstats** (JSON/CEE) every 60s to `/var/log/rsyslog_stats.json`. Typical usage:

- **Quick view (latest 20 lines):**

```
tail -n 20 /var/log/rsyslog_stats.json
```

- **Follow in real time:**

```
tail -f /var/log/rsyslog_stats.json
```

- **Parse JSON (strip the @cee: prefix) and summarize the action/queue health** — requires **jq**:

```
sed 's/^@cee: //' /var/log/rsyslog_stats.json | jq -r
'select(.name=="action" and .actionName=="lp_tls_rr") |
  "\(.timegenerated) submitted=\(.submitted) failed=\(.failed)
  suspended=\(.suspended) queuesize=\(.queuesize)'"
```

### What to watch

- **submitted** steadily increasing; **failed** near zero in normal conditions.
- **suspended=true** indicates the action is paused (e.g., all targets unavailable).
- **queuesize** and on-disk spool count/size (`du -sh /var/spool/rsyslog`) should **stabilize** then **drain** after recovery.
- For the fallback file, check growth of `/var/log/esa_fallback-buffer.log` during suspension.

### Housekeeping

- Cap disk usage via `queue.maxdiskspace` (here **10g**). Increase if necessary for longer outages.
- Review `/var/log/messages` (or `journald`) for rsyslog notices about suspension/resume events.

## 8) Rollback / change control

- To **pause LB** quickly, comment the `target=[...]` action and point to a single backend.
- To **revert packages**: disable Adiscon repo and `dnf downgrade` to the previous rsyslog build.
- Configuration is self-contained under `/etc/rsyslog*`; keep a backup before changes.

## 9) Appendix — Quick FAQ

### Q: Why not UDP for load-balancing?

A: The target pool is applied for TCP/TLS; UDP sends to the first target only. Use TCP/TLS for LB and reliability.

**Q: How does rsyslog detect target availability?**

A: Via the TCP/TLS connection lifecycle and write results. On connect or send failure, the current target is skipped; if all fail, the action is suspended and the queue buffers data until recovery.

**Q: Can we require client certs (mTLS)?**

A: Yes. Switch `StreamDriverAuthMode` from `anon` to `x509/name` and manage trust anchors accordingly.

**Q: What about TLS ciphers?**

A: Default OpenSSL policy on RHEL 9 is sane; strict ciphers can be pinned via `gtls` driver options if a policy mandates it.